

# Treatment\_\_Costs\_\_Prediction

November 12, 2023

## 0.1 Introduction

Linear regression is one of the most important algorithms under the supervised learning category in Machine Learning. It is also the commonly used model for predictive analysis. This project using this machine learning method to explore the personal health dataset and predict treatment and insurance costs.

## 0.2 Model Implementation

### 0.2.1 1. Import Data

```
[3]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[4]: # Read the data
df=pd.read_csv("insurance.csv")

# Browse the sample data
df.head()
```

```
[4]:   age    sex    bmi  children smoker    region    charges
0   19  female  27.900         0     yes southwest  16884.92400
1   18   male  33.770         1     no  southeast   1725.55230
2   28   male  33.000         3     no  southeast   4449.46200
3   33   male  22.705         0     no  northwest  21984.47061
4   32   male  28.880         0     no  northwest   3866.85520
```

### 0.2.2 2. Preprocessing the data

```
[5]: # Check for overall data information include the data types & nulls
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
```

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	age	1338 non-null	int64
1	sex	1338 non-null	object
2	bmi	1338 non-null	float64
3	children	1338 non-null	int64
4	smoker	1338 non-null	object
5	region	1338 non-null	object
6	charges	1338 non-null	float64

dtypes: float64(2), int64(2), object(3)  
memory usage: 73.3+ KB

```
[6]: # double check the NULL
df.isnull().sum()
```

```
[6]: age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

```
[7]: # Calculating some statistical data
df.describe()
```

```
[7]:
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
[8]: # Getting the data columns
df.describe().columns
```

```
[8]: Index(['age', 'bmi', 'children', 'charges'], dtype='object')
```

```
[9]: # Separating numerical and categorical data
df_num = df[['age', 'bmi', 'children', 'charges']]
df_cat = df[['sex', 'smoker', 'region']]
```

```
[10]: # Use one hot encoding to converting the categorical data into numeric data
df1 = pd.get_dummies(df_cat)
df1
```

```
[10]:
```

	sex_female	sex_male	smoker_no	smoker_yes	region_northeast	\
0	True	False	False	True	False	
1	False	True	True	False	False	
2	False	True	True	False	False	
3	False	True	True	False	False	
4	False	True	True	False	False	
...	...	...	...	...	...	
1333	False	True	True	False	False	
1334	True	False	True	False	True	
1335	True	False	True	False	False	
1336	True	False	True	False	False	
1337	True	False	False	True	False	

	region_northwest	region_southeast	region_southwest
0	False	False	True
1	False	True	False
2	False	True	False
3	True	False	False
4	True	False	False
...	...	...	...
1333	True	False	False
1334	False	False	False
1335	False	True	False
1336	False	False	True
1337	True	False	False

[1338 rows x 8 columns]

```
[11]: # Concatenating the encoded categorical and numerical data to form the dataset.
data = pd.concat([df_num,df1], axis=1)
data
```

```
[11]:
```

	age	bmi	children	charges	sex_female	sex_male	smoker_no	\
0	19	27.900	0	16884.92400	True	False	False	
1	18	33.770	1	1725.55230	False	True	True	
2	28	33.000	3	4449.46200	False	True	True	
3	33	22.705	0	21984.47061	False	True	True	
4	32	28.880	0	3866.85520	False	True	True	
...	...	...	...	...	...	...	...	
1333	50	30.970	3	10600.54830	False	True	True	
1334	18	31.920	0	2205.98080	True	False	True	
1335	18	36.850	0	1629.83350	True	False	True	
1336	21	25.800	0	2007.94500	True	False	True	

```

1337    61  29.070          0  29141.36030          True          False          False

      smoker_yes  region_northeast  region_northwest  region_southeast  \
0             True             False             False             False
1             False             False             False             True
2             False             False             False             True
3             False             False             True             False
4             False             False             True             False
...           ...               ...               ...               ...
1333          False             False             True             False
1334          False             True             False             False
1335          False             False             False             True
1336          False             False             False             False
1337          True             False             True             False

      region_southwest
0             True
1             False
2             False
3             False
4             False
...           ...
1333          False
1334          False
1335          False
1336          True
1337          False

```

[1338 rows x 12 columns]

### 0.2.3 3. Exploratory Data Analysis

```

[12]: # The correlation between the features
data.corr()

```

```

[12]:
      age      bmi  children  charges  sex_female  \
age      1.000000  0.109272  0.042469  0.299008   0.020856
bmi      0.109272  1.000000  0.012759  0.198341  -0.046371
children 0.042469  0.012759  1.000000  0.067998  -0.017163
charges  0.299008  0.198341  0.067998  1.000000  -0.057292
sex_female 0.020856 -0.046371 -0.017163 -0.057292  1.000000
sex_male  -0.020856  0.046371  0.017163  0.057292 -1.000000
smoker_no  0.025019 -0.003750 -0.007673 -0.787251  0.076185
smoker_yes -0.025019  0.003750  0.007673  0.787251 -0.076185
region_northeast 0.002475 -0.138156 -0.022808  0.006349  0.002425
region_northwest -0.000407 -0.135996  0.024806 -0.039905  0.011156

```

```

region_southeast -0.011642  0.270025 -0.023066  0.073982  -0.017117
region_southwest  0.010016 -0.006205  0.021914 -0.043210   0.004184

```

```

          sex_male  smoker_no  smoker_yes  region_northeast \
age          -0.020856   0.025019   -0.025019           0.002475
bmi           0.046371  -0.003750    0.003750          -0.138156
children      0.017163  -0.007673    0.007673          -0.022808
charges       0.057292  -0.787251    0.787251           0.006349
sex_female    -1.000000   0.076185   -0.076185           0.002425
sex_male       1.000000  -0.076185    0.076185          -0.002425
smoker_no     -0.076185   1.000000   -1.000000          -0.002811
smoker_yes     0.076185  -1.000000    1.000000           0.002811
region_northeast -0.002425  -0.002811    0.002811           1.000000
region_northwest -0.011156   0.036945   -0.036945          -0.320177
region_southeast  0.017117  -0.068498    0.068498          -0.345561
region_southwest -0.004184   0.036945   -0.036945          -0.320177

```

```

          region_northwest  region_southeast  region_southwest
age          -0.000407          -0.011642           0.010016
bmi          -0.135996           0.270025          -0.006205
children      0.024806          -0.023066           0.021914
charges       -0.039905           0.073982          -0.043210
sex_female     0.011156          -0.017117           0.004184
sex_male      -0.011156           0.017117          -0.004184
smoker_no      0.036945          -0.068498           0.036945
smoker_yes    -0.036945           0.068498          -0.036945
region_northeast -0.320177          -0.345561          -0.320177
region_northwest  1.000000          -0.346265          -0.320829
region_southeast -0.346265           1.000000          -0.346265
region_southwest -0.320829          -0.346265           1.000000

```

```

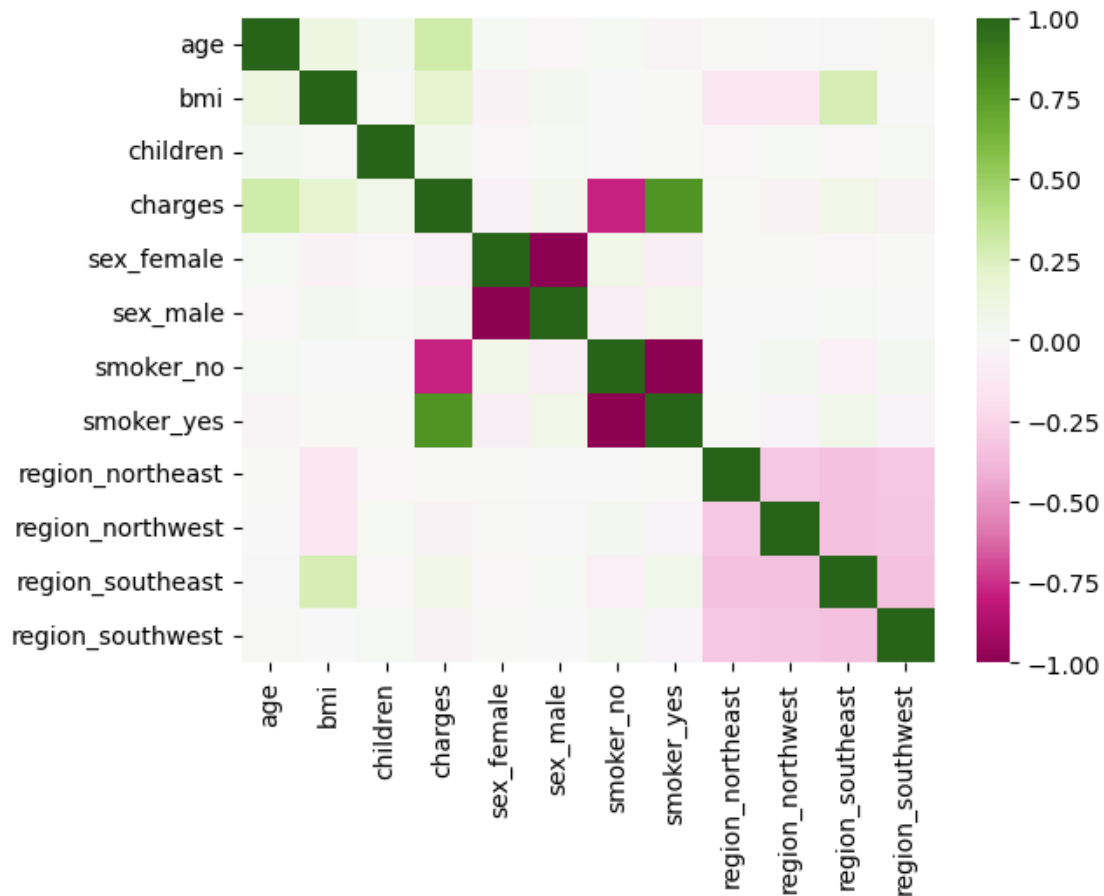
[13]: # Heatmap to visualize the correlation
      sns.heatmap(data.corr(), cmap='PiYG')

```

```

[13]: <Axes: >

```



From this heatmap we find the following observations:

1. Strong correlation between charges and smoker\_yes.
2. Weak correlation between charges and age.
3. Weak correlation between charges and bmi.
4. Weak correlation between bmi and region\_southeast.

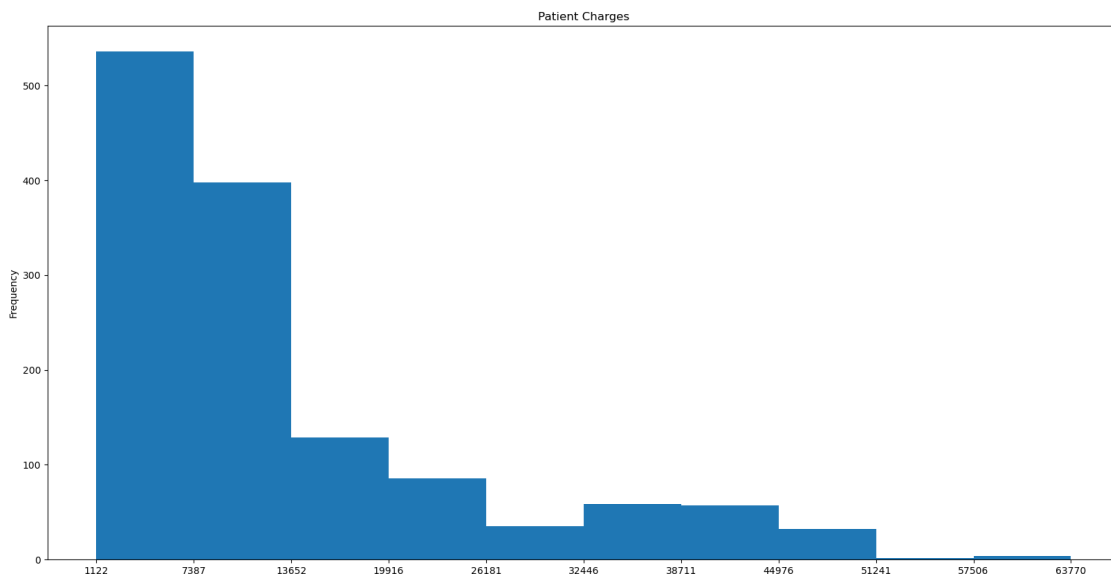
Since the values for the weak correlations are less than 0.5 so we term them as insignificant and drop them

```
[14]: # Correlation between charges and the other features.
data.corr()['charges'].sort_values()
```

```
[14]: smoker_no      -0.787251
sex_female      -0.057292
region_southwest -0.043210
region_northwest -0.039905
region_northeast  0.006349
sex_male         0.057292
children         0.067998
```

```
region_southeast    0.073982
bmi                 0.198341
age                 0.299008
smoker_yes          0.787251
charges             1.000000
Name: charges, dtype: float64
```

```
[15]: # Graph showing the min and maximum charges
count, bin_edges = np.histogram(data['charges'])
data['charges'].plot(kind='hist', xticks=bin_edges, figsize=(20,10))
plt.title("Patient Charges")
plt.show()
```



#### 0.2.4 4. Model Building

Use sklearn package to split the test and train data then use statsmodels to build a simple linear regression to predict insurance charges with the help of the other features.

```
[16]: from sklearn.model_selection import train_test_split

import statsmodels.api as sm
from statsmodels.formula.api import ols
```

### 0.2.5 5. Model fitting

For this model, we split the dataset into training and test set. We use 30% of the dataset for testing (test\_size=0.3) and then take the dataset without the charges column as the predictor variables and the charges as response/target variable.

```
[25]: x = data.drop(['charges'], axis = 1)
      y = data['charges']

      x_train,x_test,y_train,y_test = train_test_split(x,y, test_size=0.3,
      ↪random_state = 0)
      #lr = LinearRegression().fit(x_train,y_train)

      boolean_columns = ['sex_female', 'sex_male', 'smoker_no', 'smoker_yes',
      ↪'region_northeast', 'region_northwest', 'region_southeast',
      ↪'region_southwest']
      x_train[boolean_columns] = x_train[boolean_columns].astype(int)
      x_test[boolean_columns] = x_test[boolean_columns].astype(int)

[36]: # Fit the OLS model using the training data
      model_tr = sm.OLS(y_train, x_train)
      results_tr = model_tr.fit()
```

### 0.2.6 6. Model prediction

```
[39]: # Make predictions on the test set
      y_pred = results_tr.predict(x_test)

      print(r2_score(y_test,y_pred))
```

0.7909160991789904

Looks like the basic linear regression model predicting the cost of treatment look good and the score value is 0.79.

## 0.3 Model Evaluation

```
[33]: from sklearn.metrics import r2_score,mean_squared_error
```

### 0.3.1 1. Statistical Analysis:

From the summary table below, we can see that the F-statistic is 314.8 which means there are strong evidence that at least one of the independent variables in this model is related to the dependent variable “charges.” And the p-value associated with the F-statistic is very close to zero (3.47e-258) suggests that the overall model is statistically significant. While the overall model might be



significant, it's also essential to look at the significance of individual coefficients for a more detailed understanding of the contribution of each variable. In this summary table, the statistical significance of each coefficient is indicated by the “P>|t|” column. This column represents the p-value associated with the t-test for each coefficient. Therefore, the age, bmi, smoker\_no, smoker\_yes, and children are statistically significant variables for this liner regression model.

```
[38]: # Print the summary of the train OLS regression
print(results_tr.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                charges    R-squared:                0.731
Model:                        OLS        Adj. R-squared:           0.729
Method:                       Least Squares    F-statistic:             314.8
Date:                         Sun, 12 Nov 2023    Prob (F-statistic):      3.47e-258
Time:                         18:26:10        Log-Likelihood:          -9495.3
No. Observations:             936            AIC:                    1.901e+04
Df Residuals:                 927            BIC:                    1.905e+04
Df Model:                     8
Covariance Type:              nonrobust
=====
=====
=====
coef      std err          t      P>|t|      [0.025
0.975]
-----
----
age                256.4354      14.628      17.531      0.000      227.728
285.143
bmi                335.3691      34.437       9.738      0.000      267.785
402.954
children           472.7098     168.734       2.802      0.005      141.565
803.854
sex_female        -268.2715     504.493      -0.532      0.595     -1258.353
721.810
sex_male          -315.8182     509.105      -0.620      0.535     -1314.950
683.314
smoker_no         -1.201e+04     514.978     -23.321      0.000      -1.3e+04
-1.1e+04
smoker_yes         1.143e+04     543.118      21.038      0.000      1.04e+04
1.25e+04
region_northeast   443.0023     411.075       1.078      0.281      -363.742
1249.747
region_northwest  -118.8989     406.077      -0.293      0.770      -915.835
678.037
region_southeast  -551.7464     467.462      -1.180      0.238     -1469.153
365.661
region_southwest  -356.4467     419.488      -0.850      0.396     -1179.702
466.809

```

```
=====
Omnibus:                232.849    Durbin-Watson:                2.048
Prob(Omnibus):          0.000    Jarque-Bera (JB):            571.734
Skew:                   1.309    Prob(JB):                    7.07e-125
Kurtosis:               5.794    Cond. No.                     4.39e+17
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.29e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

### 0.3.2 2. the model performance

From the results below, the MSE means the squared difference between predicted and actual values is approximately 33,342,497.83 and the R-squared for test data means that the accuracy of our model is around 80% on the test data. Overall, the model is performing reasonably well on the test set, as indicated by the relatively low MSE and the high R-squared value. This is enough to conclude our model is appropriate to predict patient charges based on their personal health data

```
[50]: # Evaluate the model performance, using metrics like Mean Squared Error (MSE)
      ↪ and R2
mse = mean_squared_error(y_test, y_pred)
#mse2 = mean_squared_error(y_train,y_pred)
print(f'Mean Squared Error on Test Set: {mse}')
print(f'R2 for test data: {r2_score(y_test,y_pred)}')
```

Mean Squared Error on Test Set: 33342497.82695458

R2 for test data: 0.7909160991789904

## 0.4 References

1. Miri Choi, Medical Cost Personal Datasets (2013), Kaggle, <https://www.kaggle.com/datasets/mirichoi0218/insurance>
2. Thomas George, Predicting Patient treatment costs (2020), Medium, <https://medium.com/analytics-vidhya/>

```
[ ]:
```